

END-TO-END ML PIPELINE WITH DEPLOYMENT

1. Project Overview

This project focuses on building an end-to-end **Machine Learning Pipeline** capable of fully automated data ingestion, feature engineering, model training, hyperparameter tuning, evaluation, and deployment.

The system is designed to be scalable, modular, and production-ready, with CI/CD-driven model retraining and automatic artifact versioning.

The pipeline supports multiple ML algorithms—including **XGBoost**, **LightGBM**, and **Random Forest**—and selects the best-performing model based on validation metrics.

The final model is deployed as a **REST API** using FastAPI/Flask, enabling real-time predictions.

2. Problem Statement

Organizations often struggle with ML models that:

- Break due to data drift
- Are trained manually without automation
- Lack reproducibility
- Cannot be deployed easily
- Perform inconsistently across datasets

This project solves these challenges by building a **production-grade ML workflow** that automates the entire lifecycle from data to deployment.

3. Objectives

- Build a modular ML pipeline with reusable components
- Automate ingestion, cleaning, and feature extraction
- Integrate model training with hyperparameter tuning

- Support multiple algorithms with automated model selection
 - Deploy the best model as an inference API
 - Enable reproducible experiments using proper artifact logging
 - Add CI/CD automation for scheduled retraining
-

4. Tech Stack

Languages

- Python

Frameworks & Libraries

- pandas, NumPy
- scikit-learn
- XGBoost, LightGBM
- joblib, pickle
- FastAPI / Flask
- Docker
- GitHub Actions (CI/CD)

Tools

- Git
- DockerHub
- VS Code
- Postman for API testing

5. System Architecture



6. Data Pipeline

6.1 Ingestion

- Loaded raw data from CSV files
- Validated schema (rows, missing values, data types)
- Added timestamp-based indexing

6.2 Cleaning

- Imputed missing values using median strategy
- Removed outliers using IQR
- Normalized skewed distributions
- Checked data leakage

6.3 Feature Engineering

- Scaled numerical features with StandardScaler
- Encoded categorical variables
- Generated interaction features
- Removed low-variance features
- Applied correlation-based feature selection

6.4 Dataset Splitting

A 70–10–20 split was used:

- 70% training
 - 10% validation
 - 20% testing
-

7. Model Development

Models Used

- **Random Forest Classifier**
- **XGBoost Classifier**
- **LightGBM Classifier**

Each model underwent:

- **Hyperparameter tuning via GridSearchCV**
- **Cross-validation**
- **Class imbalance handling**
- **Performance scoring on validation data**

Model Selection Criteria

The best model was chosen using:

- **Accuracy**
- **F1-score**
- **ROC-AUC**
- **Validation loss**

The pipeline automatically logs:

- **Best model**
 - **Metrics**
 - **Hyperparameters**
-

8. Evaluation Metrics

Model	Accuracy	F1 Score	AUC
Random Forest	0.91	0.89	0.93
XGBoost	0.92	0.90	0.94
LightGBM	0.93	0.92	0.96

LightGBM was selected as the final deployed model.

9. Deployment

9.1 API Development

A FastAPI/Flask service was created to serve real-time predictions.

Routes:

- `/predict` — returns prediction + probabilities
- `/health` — health check endpoint

9.2 Dockerization

- Containerized app using Docker
- Pushed to DockerHub
- Prepared image for cloud deployment

9.3 CI/CD (Retraining + Deployment)

Using GitHub Actions:

- On data update → trigger retraining
- On code update → run unit tests
- Auto-build & deploy Docker image

- Notify via GitHub checks
-

10. Challenges & Solutions

Challenge 1: Data Drift Across Time

Solution: Added drift detection module + scheduled retraining.

Challenge 2: Large Feature Space

Solution: Automated feature importance pruning based on SHAP + permutation importance.

Challenge 3: Maintaining Reproducibility

Solution:

- Versioning models
 - Logging dependencies in requirements.txt
 - Saving scaler + encoder for future predictions
-

11. Conclusion

The project successfully delivers an automated, production-ready ML workflow with:

- End-to-end pipeline
- Multiple model support
- Automated selection
- API deployment
- Drift detection and retraining

This system can be applied to fintech, healthcare, retail forecasting, fraud detection, and other industry ML use cases.

